

ECE 695 Task 2 Report

Yizhou Shan <shan13@purdue.edu>
Wan-Eih Huang <huan1031@purdue.edu>

I. Overall

In task 2, we mainly used `pte_none()` and `pte_young()` to identify if a page has been referenced. And we can only count reference to 1.

A page was referenced if and only if these two conditions are true:
1) PTE has valid bit set, which means `pte_none()` returns false, 2)
PTE has accessed bit set, which means `pte_young()` returns true.
Otherwise, this page was not referenced.

II. Detail

We changed two source files, 1) `fs/proc/base.c`, 2) `fs/proc/task_mmu.c`. Originally, the `/proc/PID/maps` is read-only, we make it writable so we are able to turn on or turn off our own feature during runtime.

By default, the `/proc/PID/maps` will print its default content.
To enable our feature, type: **`echo y > /proc/PID/maps`**.
To reset, type: **`echo n > /proc/PID/maps`**.

Our main API is:

```
show_vma_page_map(struct seq_file *m,  
                  struct vm_area_struct *vma,  
                  unsigned long start,  
                  unsigned long end)
```

which is implemented similar to `unmap_page_range()` function. Our function will walk through the user page table through standard kernel pgtable APIs, and check PTE valid and accessed bit.

III. Interesting Things

The last `[vsyscall]` segment actually does not have an associated `mm_struct`. The `vsyscall` and `vDSO` are both used to accelerate system call, to avoid frequent context switch.